

# Cloud Databases

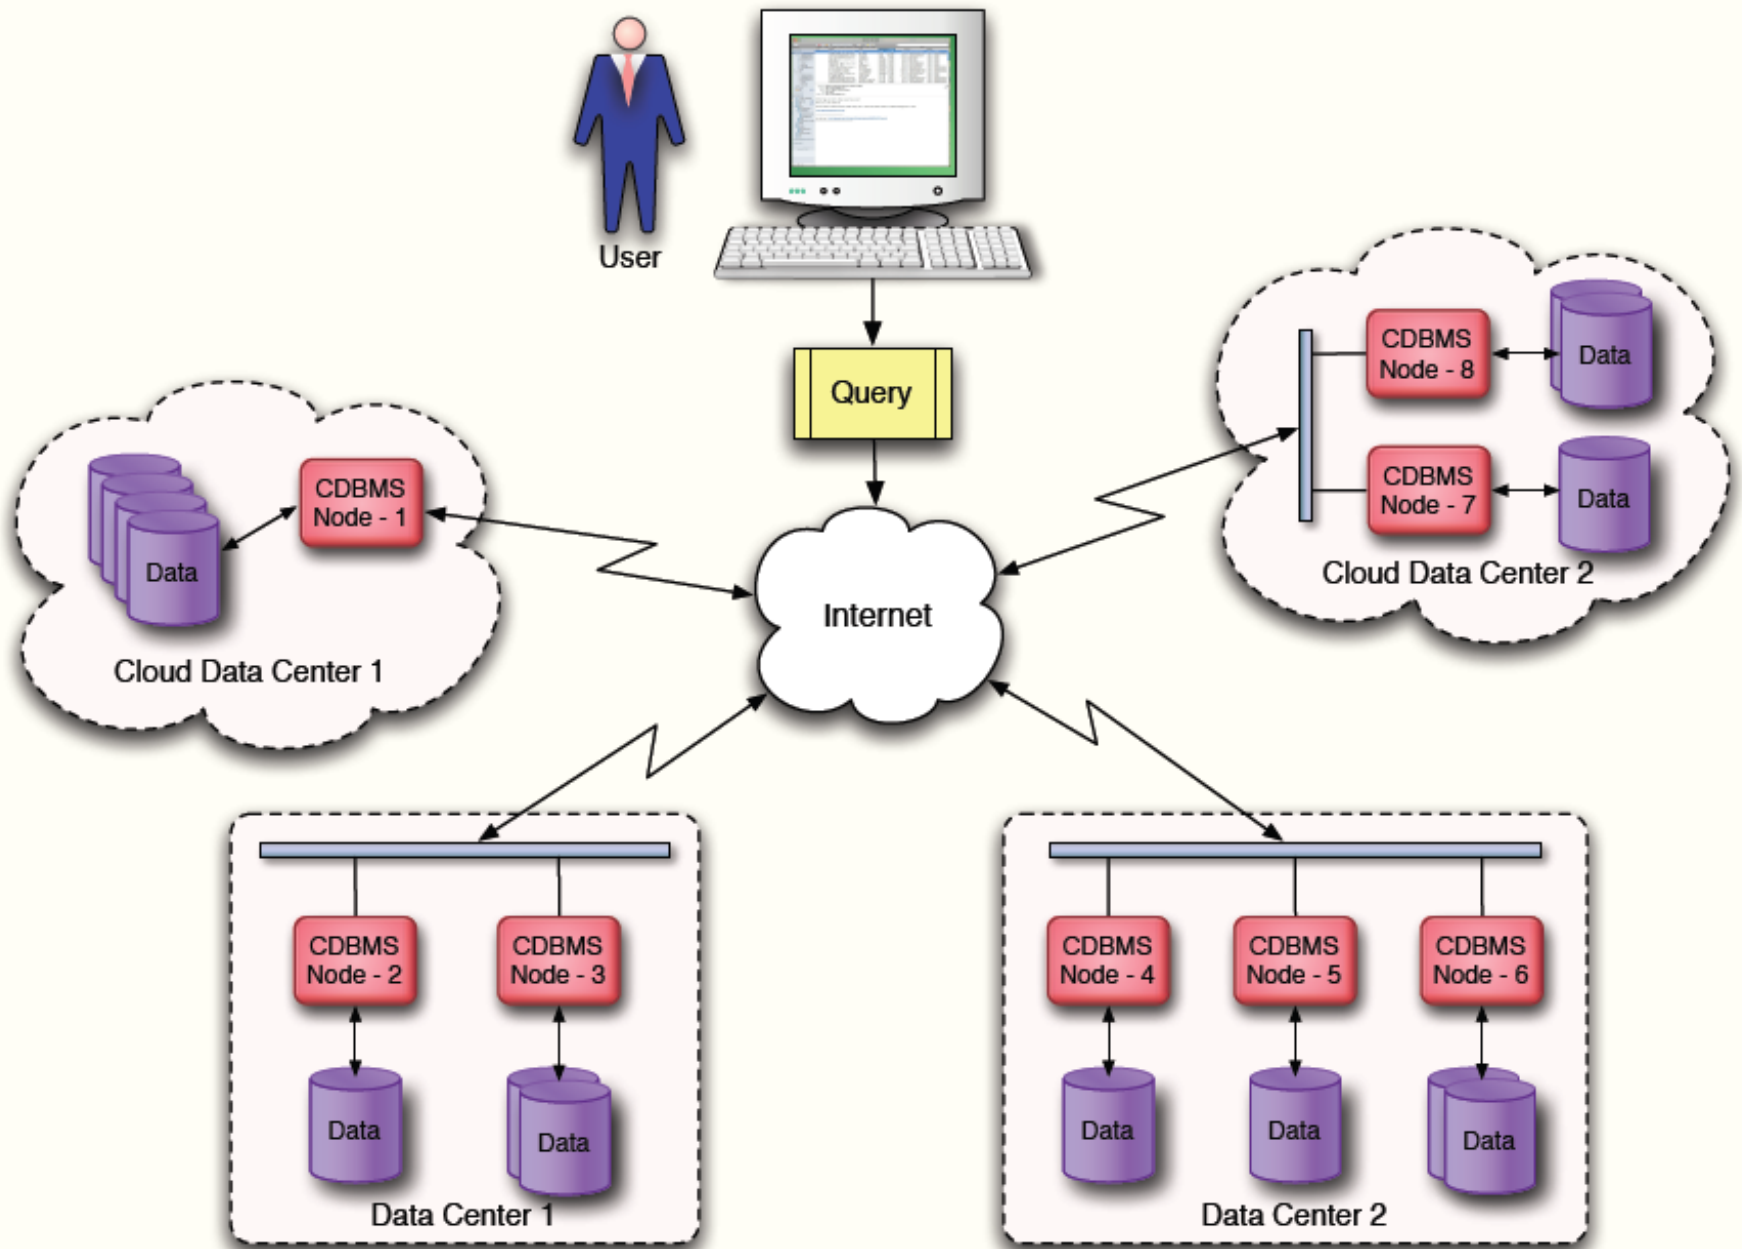
# Introduction

- Std. Layered Architecture
- **Cloud Infrastructure**
- A **cloud database** is a database that typically runs on a *cloud computing* platform, such as *Amazon EC2, GoGrid, Salesforce, Rackspace*, and *Microsoft Azure*.
- Deployment models
  - users can run databases on the cloud independently, using a *virtual machine* image
  - they can purchase access to a database service, maintained by a cloud database provider. **DBaaS**
- Data Models : **SQL , NoSQL**

# Definition

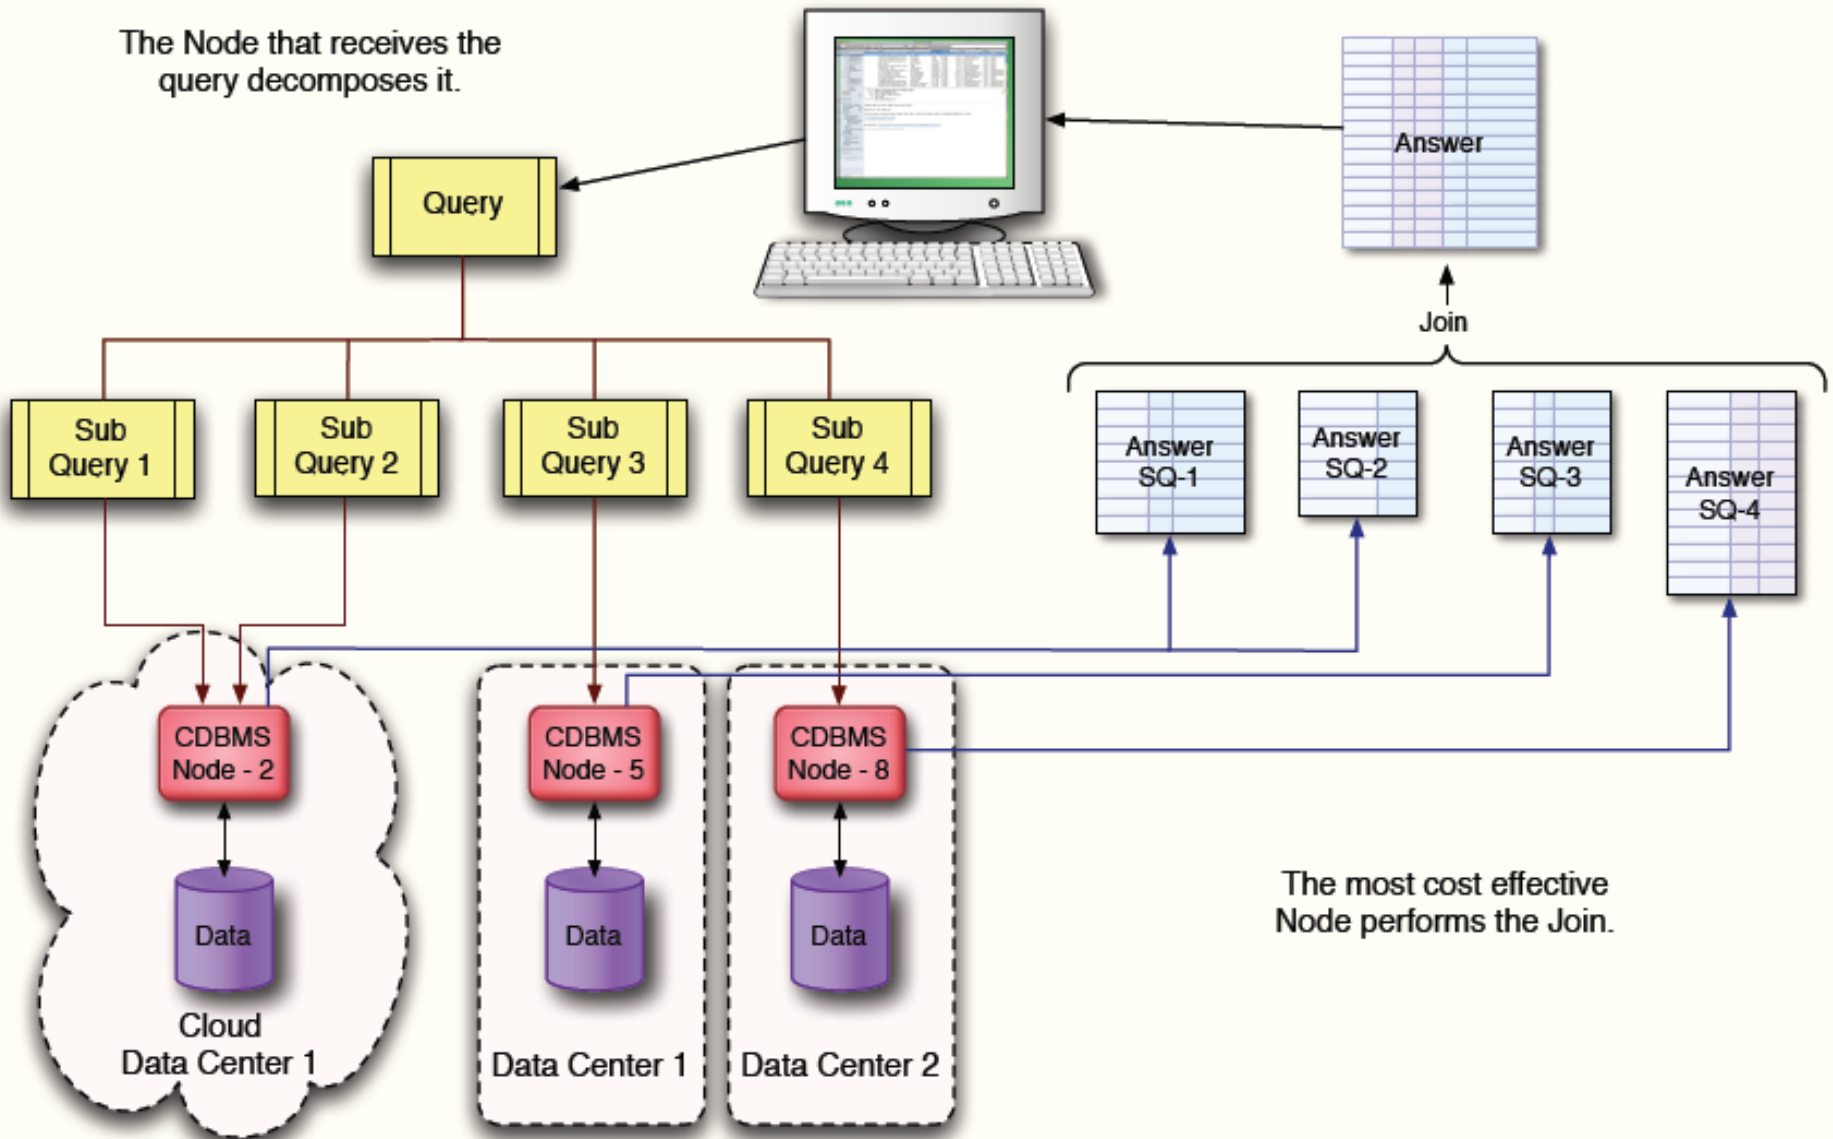
**Cloud dbms (CDBMS) is a distributed database that delivers a query service across multiple distributed database nodes located in multiple geographically-distributed data centers, both corporate data centers and cloud data centers.**

# Architecture / Layout



# Distributed Query Processing

The Node that receives the query decomposes it.



The most cost effective Node performs the Join.



# Data Model : SQL

- **SQL database**, such as *NuoDB*, *Oracle Database*, *Microsoft SQL Server*, and *MySQL*, are one type of database which can be run on the cloud (either as a Virtual Machine Image or as a service, depending on the vendor).
- SQL databases are difficult to scale, not natively suited to a cloud environment
- Cloud database services based on SQL are attempting to address this challenge

# Data Model : NoSQL [www.nosql-database.org](http://www.nosql-database.org)

- NoSQL means 'Not Only SQL', 'Not Relational'.
- NoSQL databases, such as *Apache Cassandra*, *CouchDB* and *MongoDB*, are another type of database which can run on the cloud.
- NoSQL databases are built to service heavy read/write loads and are able scale up and down easily
- More natively suited to running on the cloud.
- working with NoSQL databases often requires a complete rewrite of application code
- Set of APIs to access data. *no SQL like query*

# NoSQL : Advantages

- non-relational
- don't require schema
- data are replicated to multiple nodes (so, identical & fault-tolerant) and can be partitioned:
  - down nodes easily replaced
  - no single point of failure
- horizontal scalable
- cheap, easy to implement (open-source)
- massive write performance
- fast key-value access





# NoSQL : Disadvantages

- Don't fully support relational features
  - no join, group by, order by operations (except within partitions)
  - no referential integrity constraints across partitions
- No declarative query language (e.g., SQL) → more programming
- Relaxed ACID (see CAP theorem) → fewer guarantees
- No easy integration with other applications that support SQL

# NOSQL Modeling Types

## 1.Key-value

- Example: DynamoDB, Voldermort, Scalaris

## 2.Document-based

- Example: MongoDB, CouchDB

## 3.Column-based

- Example: BigTable, Cassandra, Hbased

## 4.Graph-based

- Example: Neo4J, InfoGrid
- “No-schema” is a common characteristics of most NOSQL storage systems
- Provide “flexible” data types

# NoSQL Transactions

## Types of consistency:

1. Strong consistency – ACID  
(**A**tomicity, **C**onsistency, **I**solation, **D**urability)  
*do not supported by NoSQL*
2. Weak consistency – BASE  
(**B**asically **A**vailable **S**oft- state **E**ventual consistency)

***Based on CAP Theorem***

# CAP Theorem

- Three properties of a distributed system (sharing data)

- **Consistency:**


- all copies have same value

- **Availability:**

- reads and writes always succeed

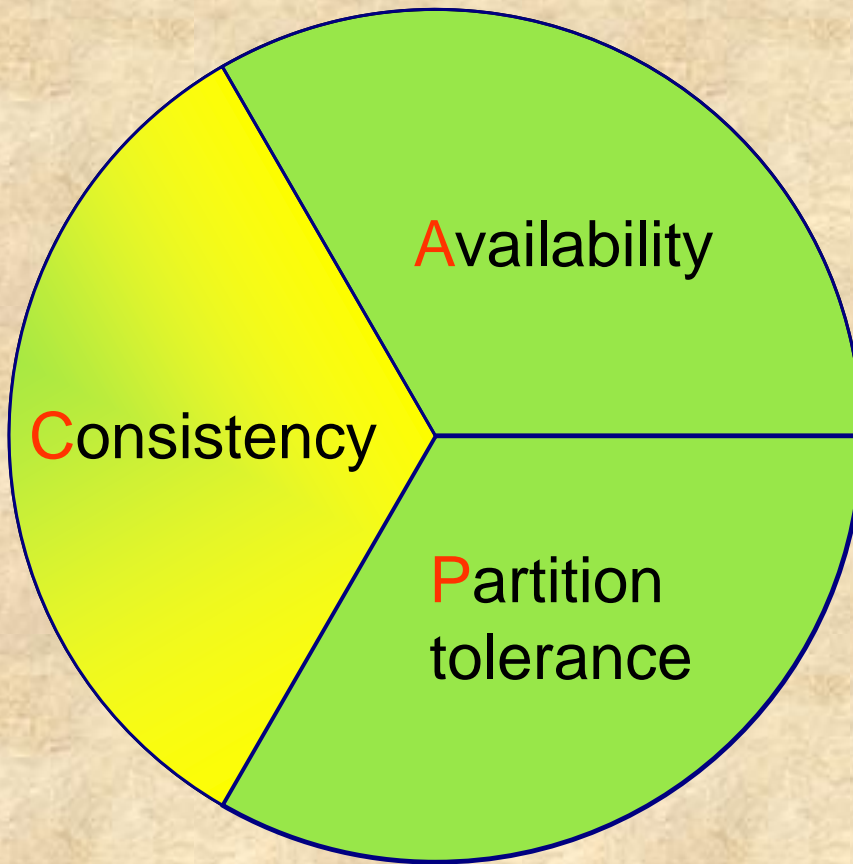
- **Partition-tolerance:**

- system properties (consistency and/or availability) hold even when network failures prevent some machines from communicating with others



# CAP Theorem

All client always have the same view of the data

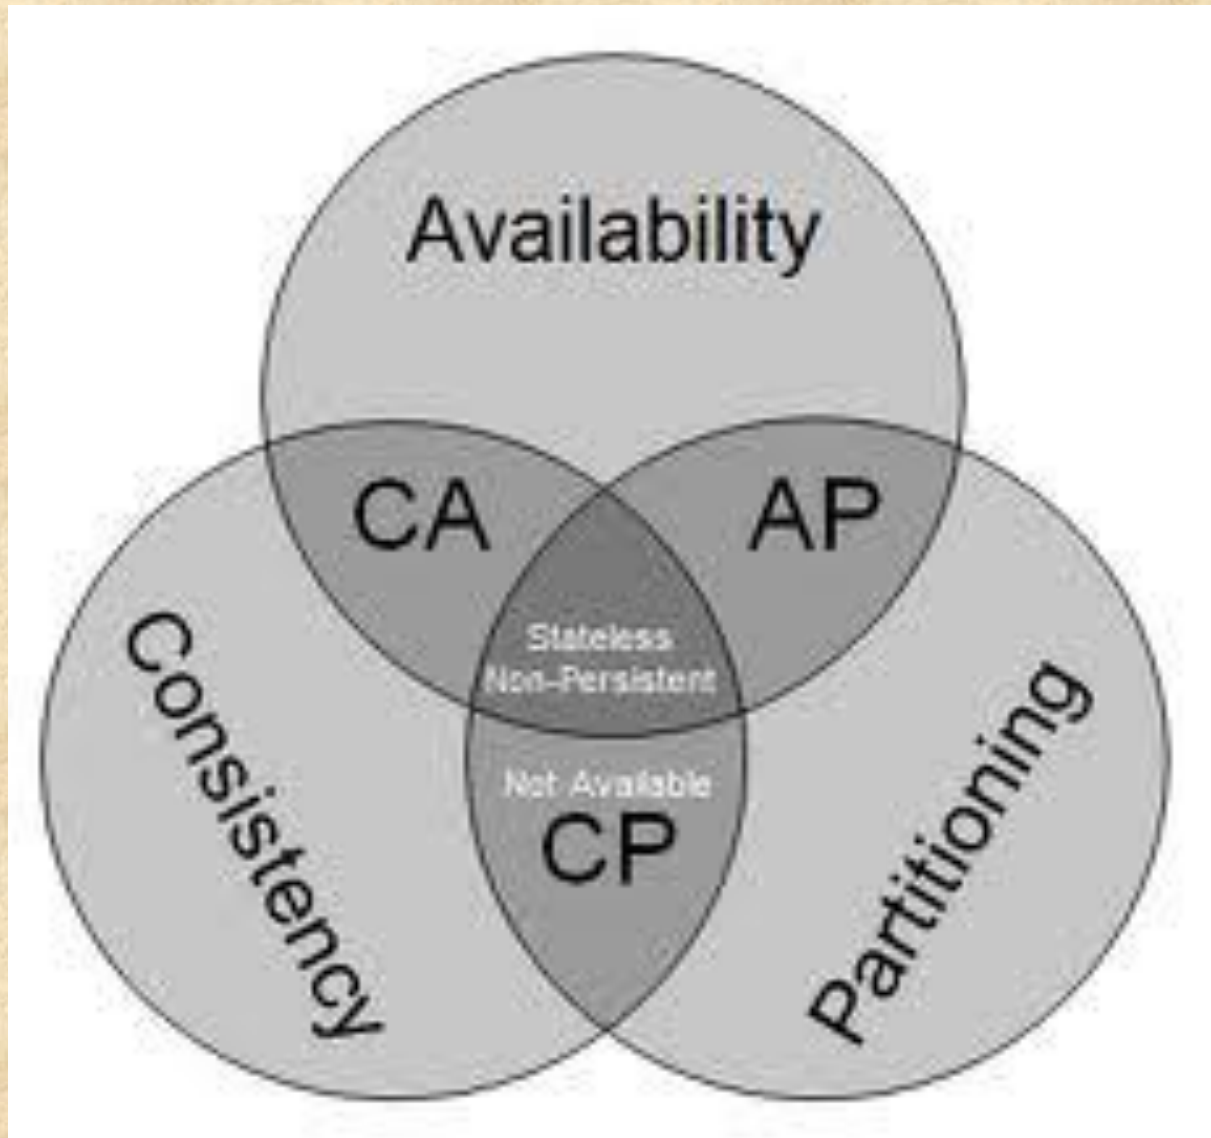




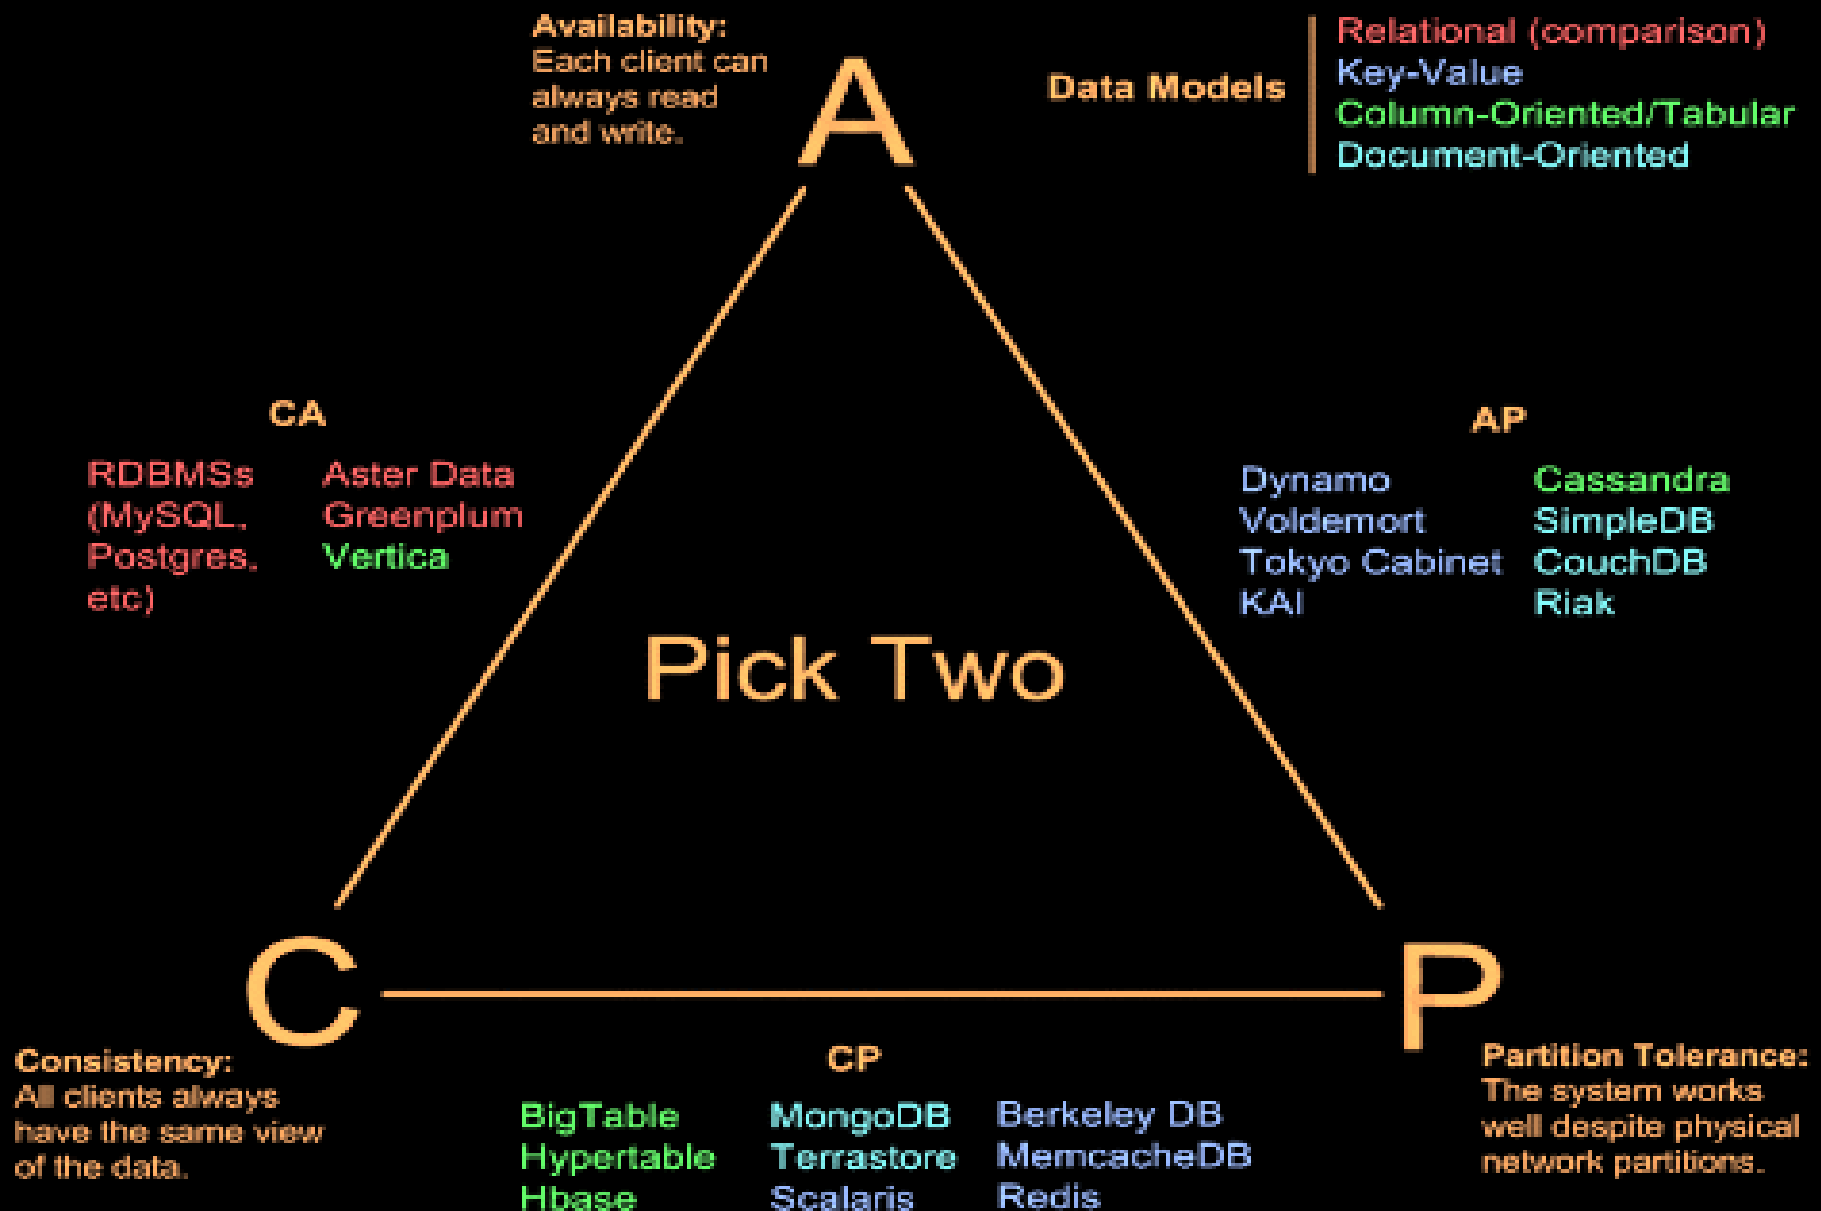
# Brewer's CAP Theorem:

- *For any system sharing data, it is “impossible” to guarantee simultaneously all of these three properties*
- You can have at most two of these three properties for any shared-data system
- Very large systems will “partition” at some point:
  - That leaves either **C** or **A** to choose from (traditional DBMS prefers **C** over **A** and **P** )
  - In almost all cases, you would choose **A** over **C** (except in specific applications such as order processing)

# Brewer's CAP Theorem



# Visual Guide to NoSQL Systems



# Storage Architecture for Cloud DB

## Shared-nothing Storage Architecture

- It involves data partitioning which splits the data into independent sets - physically located on different database servers.
- suitable for Cloud.
- Needs piece of middleware to route database requests to the appropriate server.
- IBM, Oracle, Amazon's SimpleDB, Hadoop Distributed File System and Yahoo's PNUTS also implement shared-nothing architecture

# Storage Architecture for Cloud DB

## Shared-disk Database Architecture

- Treats the whole database as a single large piece of database stored on a Storage Area Network (SAN) or Network Attached Storage (NAS) storage that is shared and accessible through network by all nodes.
- Middleware is not required to route data requests to specific servers as each node/client has access to all of the data.
- Oracle RAC, IBM DB2 pureScale, Sybase etc. support this architecture



# Apache Cassandra

# CouchDB

# MongoDB

# Comparison of RDBMS and NoSQL databases

RDBMS	NoSQL Databases
<ul style="list-style-type: none"><li>• Data within a database is treated as a “whole”</li></ul>	<ul style="list-style-type: none"><li>• Each entity is considered an independent unit of data and can be freely moved from one machine to the other</li></ul>
<ul style="list-style-type: none"><li>• RDBMS support centrally managed architecture.</li></ul>	<ul style="list-style-type: none"><li>• They follow distributed architecture.</li></ul>
<ul style="list-style-type: none"><li>• They are statically provisioned.</li></ul>	<ul style="list-style-type: none"><li>• They are dynamically provisioned.</li></ul>
<ul style="list-style-type: none"><li>• It is difficult to scale them.</li></ul>	<ul style="list-style-type: none"><li>• They are easily scalable.</li></ul>
<ul style="list-style-type: none"><li>• They provide SQL to query data</li></ul>	<ul style="list-style-type: none"><li>• They use API to query data (not feature rich as SQL).</li></ul>
<ul style="list-style-type: none"><li>• ACID (Atomicity, Consistency, Isolation and Durability) Compliant; DBMS maintains Consistency.</li></ul>	<ul style="list-style-type: none"><li>• Follow BASE (Basically Available, Soft state, Eventually consistent); The user accesses are guaranteed only at a single-key level.</li></ul>
<ul style="list-style-type: none"><li>• They support on-line Transaction Processing applications.</li></ul>	<ul style="list-style-type: none"><li>• They support web2.0 applications.</li></ul>

# Challenges to Develop Cloud Databases

